

MNIST_GAN_Exercise

March 2, 2019

1 Generative Adversarial Network

In this notebook, we'll be building a generative adversarial network (GAN) trained on the MNIST dataset. From this, we'll be able to generate new handwritten digits!

GANs were [first reported on](#) in 2014 from Ian Goodfellow and others in Yoshua Bengio's lab. Since then, GANs have exploded in popularity. Here are a few examples to check out:

- [Pix2Pix](#)
- [CycleGAN & Pix2Pix in PyTorch, Jun-Yan Zhu](#)
- [A list of generative models](#)

The idea behind GANs is that you have two networks, a generator G and a discriminator D , competing against each other. The generator makes "fake" data to pass to the discriminator. The discriminator also sees real training data and predicts if the data it's received is real or fake. > * The generator is trained to fool the discriminator, it wants to output data that looks *as close as possible* to real, training data. * The discriminator is a classifier that is trained to figure out which data is real and which is fake.

What ends up happening is that the generator learns to make data that is indistinguishable from real data to the discriminator.

The general structure of a GAN is shown in the diagram above, using MNIST images as data. The latent sample is a random vector that the generator uses to construct its fake images. This is often called a **latent vector** and that vector space is called **latent space**. As the generator trains, it figures out how to map latent vectors to recognizable images that can fool the discriminator.

If you're interested in generating only new images, you can throw out the discriminator after training. In this notebook, I'll show you how to define and train these adversarial networks in PyTorch and generate new images!

```
In [27]: %matplotlib inline
```

```
import numpy as np
import torch
import matplotlib.pyplot as plt
```

```
In [28]: from torchvision import datasets
import torchvision.transforms as transforms
```

```
# number of subprocesses to use for data loading
num_workers = 0
```

```

# how many samples per batch to load
batch_size = 64

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# get the training datasets
train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transform)

# prepare data loader
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers)

```

1.0.1 Visualize the data

```

In [29]: # obtain one batch of training images
dataiter = iter(train_loader)
images, labels = dataiter.next()
images = images.numpy()

```

```

# get one image from the batch
img = np.squeeze(images[0])


fig = plt.figure(figsize = (3,3))
ax = fig.add_subplot(111)
ax.imshow(img, cmap='gray')

```

```

Out[29]: <matplotlib.image.AxesImage at 0x7fa4cafd4a90>

```



2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

2.1 Discriminator

The discriminator network is going to be a pretty typical linear classifier. To make this network a universal function approximator, we'll need at least one hidden layer, and these hidden layers should have one key attribute: > All hidden layers will have a [Leaky ReLU](#) activation function applied to their outputs.

Leaky ReLU We should use a leaky ReLU to allow gradients to flow backwards through the layer unimpeded. A leaky ReLU is like a normal ReLU, except that there is a small non-zero output for negative input values.

Sigmoid Output We'll also take the approach of using a more numerically stable loss function on the outputs. Recall that we want the discriminator to output a value 0-1 indicating whether an image is *real or fake*. > We will ultimately use [BCEWithLogitsLoss](#), which combines a sigmoid activation function **and** binary cross entropy loss in one function.

So, our final output layer should not have any activation function applied to it.

```
In [30]: import torch.nn as nn
import torch.nn.functional as F

class Discriminator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Discriminator, self).__init__()

        self.input_size = input_size

        # define all layers
        self.hidden1 = nn.Linear(input_size, 4*hidden_dim) # dimensions are reduced
        self.hidden2 = nn.Linear(4*hidden_dim, 2*hidden_dim)
        self.hidden3 = nn.Linear(2*hidden_dim, hidden_dim)
        self.output = nn.Linear(hidden_dim, output_size)

        # dropout layer
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        # flatten image
        x = x.view(-1, 28*28)

        # pass x through all layers
        # apply leaky relu activation to all hidden layers
        x = F.leaky_relu(self.hidden1(x), 0.2) # (input, negative_slope=0.2)
        x = self.dropout(x)
```

```

x = F.leaky_relu(self.hidden2(x), 0.2)
x = self.dropout(x)
x = F.leaky_relu(self.hidden3(x), 0.2)
x = self.dropout(x)
x = self.output(x) # no activation here (sigmoid will be applied later)

return x

```

2.2 Generator

The generator network will be almost exactly the same as the discriminator network, except that we're applying a [tanh activation function](#) to our output layer.

tanh Output The generator has been found to perform the best with *tanh* for the generator output, which scales the output to be between -1 and 1, instead of 0 and 1.

Recall that we also want these outputs to be comparable to the *real* input pixel values, which are read in as normalized values between 0 and 1. > So, we'll also have to **scale our real input images to have pixel values between -1 and 1** when we train the discriminator.

I'll do this in the training loop, later on.

In [31]: `class Generator(nn.Module):`

```

def __init__(self, input_size, hidden_dim, output_size):
    super(Generator, self).__init__()

    # define all layers
    self.hidden1 = nn.Linear(input_size, hidden_dim) # dims are increased (unlike d
    self.hidden2 = nn.Linear(hidden_dim, 2*hidden_dim)
    self.hidden3 = nn.Linear(2*hidden_dim, 4*hidden_dim)
    self.output = nn.Linear(4*hidden_dim, output_size)

    # dropout layer
    self.dropout = nn.Dropout(0.3)

def forward(self, x):
    # pass x through all layers

    # final layer should have tanh applied
    x = F.leaky_relu(self.hidden1(x), 0.2) # (input, negative_slope=0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.hidden2(x), 0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.hidden3(x), 0.2)
    x = self.dropout(x)
    x = F.tanh(self.output(x)) # tanh activation: scaling between -1 and 1

    return x

```

2.3 Model hyperparameters

```
In [32]: # Discriminator hyperparams

# Size of input image to discriminator (28*28)
input_size = 28*28
# Size of discriminator output (real or fake)
d_output_size = 1
# Size of *last* hidden layer in the discriminator
d_hidden_size = 28

# Generator hyperparams

# Size of latent vector to give to generator
z_size = 100
# Size of discriminator output (generated image)
g_output_size = input_size
# Size of *first* hidden layer in the generator
g_hidden_size = 28
```

2.4 Build complete network

Now we're instantiating the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

```
In [33]: # instantiate discriminator and generator
D = Discriminator(input_size, d_hidden_size, d_output_size)
G = Generator(z_size, g_hidden_size, g_output_size)

# check that they are as you expect
print(D)
print()
print(G)
```

```
Discriminator(
  (hidden1): Linear(in_features=784, out_features=112, bias=True)
  (hidden2): Linear(in_features=112, out_features=56, bias=True)
  (hidden3): Linear(in_features=56, out_features=28, bias=True)
  (output): Linear(in_features=28, out_features=1, bias=True)
  (dropout): Dropout(p=0.3)
)
```

```
Generator(
  (hidden1): Linear(in_features=100, out_features=28, bias=True)
  (hidden2): Linear(in_features=28, out_features=56, bias=True)
  (hidden3): Linear(in_features=56, out_features=112, bias=True)
  (output): Linear(in_features=112, out_features=784, bias=True)
  (dropout): Dropout(p=0.3)
```

)

2.5 Discriminator and Generator Losses

Now we need to calculate the losses.

2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images, $d_loss = d_real_loss + d_fake_loss$.
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

The losses will be binary cross entropy loss with logits, which we can get with `BCEWithLogitsLoss`. This combines a sigmoid activation function **and** binary cross entropy loss in one function.

For the real images, we want $D(\text{real_images}) = 1$. That is, we want the discriminator to classify the real images with a label = 1, indicating that these are real. To help the discriminator generalize better, the labels are **reduced a bit from 1.0 to 0.9**. For this, we'll use the parameter `smooth`; if `True`, then we should smooth our labels. In PyTorch, this looks like `labels = torch.ones(size) * 0.9`

The discriminator loss for the fake data is similar. We want $D(\text{fake_images}) = 0$, where the fake images are the *generator output*, `fake_images = G(z)`.

2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get $D(\text{fake_images}) = 1$. In this case, the labels are **flipped** to represent that the generator is trying to fool the discriminator into thinking that the images it generates (fakes) are real!

```
In [34]: # Calculate losses
def real_loss(D_out, smooth=False):
    # compare logits to real labels
    # smooth labels if smooth=True
    batch_size = D_out.size(0)
    if(smooth):
        labels = torch.ones(batch_size) * 0.9
    else:
        labels = torch.ones(batch_size)

    criterion = nn.BCEWithLogitsLoss()
    # calc loss
    loss = criterion(D_out.squeeze(), labels)

    return loss
```

```
def fake_loss(D_out):
    # compare logits to fake labels
    batch_size = D_out.size(0)
    labels = torch.zeros(batch_size)

    criterion = nn.BCEWithLogitsLoss()
    # calc loss
    loss = criterion(D_out.squeeze(), labels)
    return loss
```

2.6 Optimizers

We want to update the generator and discriminator variables separately. So, we'll define two separate Adam optimizers.

```
In [35]: import torch.optim as optim

        # learning rate for optimizers
        lr = 0.002

        # Create optimizers for the discriminator and generator
        d_optimizer = optim.Adam(D.parameters(), lr)
        g_optimizer = optim.Adam(G.parameters(), lr)
```

2.7 Training

Training will involve alternating between training the discriminator and the generator. We'll use our functions `real_loss` and `fake_loss` to help us calculate the discriminator losses in all of the following cases.

2.7.1 Discriminator training

1. Compute the discriminator loss on real, training images
2. Generate fake images
3. Compute the discriminator loss on fake, generated images
4. Add up real and fake loss
5. Perform backpropagation + an optimization step to update the discriminator's weights

2.7.2 Generator training

1. Generate fake images
2. Compute the discriminator loss on fake images, using **flipped** labels!
3. Perform backpropagation + an optimization step to update the generator's weights

Saving Samples As we train, we'll also print out some loss statistics and save some generated "fake" samples.

```
In [43]: import pickle as pkl
```

```
# training hyperparams
num_epochs = 100

# keep track of loss and generated, "fake" samples
samples = []
losses = []

print_every = 400

# Get some fixed data for sampling. These are images that are held
# constant throughout training, and allow us to inspect the model's performance
sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()

# train the network
D.train()
G.train()
for epoch in range(num_epochs):

    for batch_i, (real_images, _) in enumerate(train_loader):

        batch_size = real_images.size(0)

        ## Important rescaling step ##
        real_images = real_images*2 - 1 # rescale input images from [0,1) to [-1, 1)

        # =====
        #                TRAIN THE DISCRIMINATOR
        # =====
        d_optimizer.zero_grad()

        # 1. Train with real images

        # Compute the discriminator losses on real images
        # use smoothed labels
        D_out = D.forward(real_images)
        d_real_loss = real_loss(D_out, smooth=True)

        # 2. Train with fake images

        # Generate fake images
        z = np.random.uniform(-1, 1, size=(batch_size, z_size))
```



```

z = torch.from_numpy(z).float()
fake_images = G.forward(z)

# Compute the discriminator losses on fake images
D_out = D.forward(fake_images)
d_fake_loss = fake_loss(D_out)

# add up real and fake losses and perform backprop
d_loss = d_real_loss + d_fake_loss
d_loss.backward()
d_optimizer.step()

# =====
#                     TRAIN THE GENERATOR
# =====
g_optimizer.zero_grad()

# 1. Train with fake images and flipped labels

# Generate fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G.forward(z)

# Compute the discriminator losses on fake images
# using flipped labels!
D_out = D.forward(fake_images)
g_loss = real_loss(D_out) # adversarial: real_loss instead of fake_loss

# perform backprop
g_loss.backward()
g_optimizer.step()

# Print some loss stats
if batch_i % print_every == 0:
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.64f} | g_loss: {:.64f}'.format(
        epoch+1, num_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# append discriminator loss and generator loss
losses.append((d_loss.item(), g_loss.item()))

# generate and save sample, fake images
G.eval() # eval mode for generating samples
samples_z = G(fixed_z)

```

```

        samples.append(samples_z)
    G.train() # back to train mode

    # Save training generator samples
    with open('train_samples.pkl', 'wb') as f:
        pickle.dump(samples, f)

```

```

Epoch [ 1/ 100] | d_loss: 1.2913 | g_loss: 0.9683
Epoch [ 1/ 100] | d_loss: 1.2203 | g_loss: 1.0513
Epoch [ 1/ 100] | d_loss: 1.3839 | g_loss: 1.0690
Epoch [ 2/ 100] | d_loss: 1.1739 | g_loss: 1.1843
Epoch [ 2/ 100] | d_loss: 1.1875 | g_loss: 1.2010
Epoch [ 2/ 100] | d_loss: 1.2770 | g_loss: 1.0769
Epoch [ 3/ 100] | d_loss: 1.2870 | g_loss: 1.0315
Epoch [ 3/ 100] | d_loss: 1.2843 | g_loss: 1.1508
Epoch [ 3/ 100] | d_loss: 1.3419 | g_loss: 0.9113
Epoch [ 4/ 100] | d_loss: 1.4190 | g_loss: 0.8892
Epoch [ 4/ 100] | d_loss: 1.3236 | g_loss: 0.9185
Epoch [ 4/ 100] | d_loss: 1.2054 | g_loss: 1.0921
Epoch [ 5/ 100] | d_loss: 1.2896 | g_loss: 0.8734
Epoch [ 5/ 100] | d_loss: 1.2756 | g_loss: 0.9989
Epoch [ 5/ 100] | d_loss: 1.3833 | g_loss: 0.9196
Epoch [ 6/ 100] | d_loss: 1.2616 | g_loss: 0.9206
Epoch [ 6/ 100] | d_loss: 1.2460 | g_loss: 0.9680
Epoch [ 6/ 100] | d_loss: 1.3425 | g_loss: 0.8214
Epoch [ 7/ 100] | d_loss: 1.3468 | g_loss: 0.9200
Epoch [ 7/ 100] | d_loss: 1.3098 | g_loss: 1.0249
Epoch [ 7/ 100] | d_loss: 1.2890 | g_loss: 0.9971
Epoch [ 8/ 100] | d_loss: 1.2636 | g_loss: 0.9520
Epoch [ 8/ 100] | d_loss: 1.2242 | g_loss: 0.9332
Epoch [ 8/ 100] | d_loss: 1.3774 | g_loss: 1.0290
Epoch [ 9/ 100] | d_loss: 1.2166 | g_loss: 1.1847
Epoch [ 9/ 100] | d_loss: 1.2757 | g_loss: 0.8795
Epoch [ 9/ 100] | d_loss: 1.3350 | g_loss: 0.9848
Epoch [ 10/ 100] | d_loss: 1.2258 | g_loss: 1.2294
Epoch [ 10/ 100] | d_loss: 1.2826 | g_loss: 1.0322
Epoch [ 10/ 100] | d_loss: 1.3249 | g_loss: 0.7789
Epoch [ 11/ 100] | d_loss: 1.4087 | g_loss: 0.8040
Epoch [ 11/ 100] | d_loss: 1.3091 | g_loss: 0.9468
Epoch [ 11/ 100] | d_loss: 1.3661 | g_loss: 0.8582
Epoch [ 12/ 100] | d_loss: 1.2972 | g_loss: 1.0706
Epoch [ 12/ 100] | d_loss: 1.2910 | g_loss: 0.9599
Epoch [ 12/ 100] | d_loss: 1.4469 | g_loss: 0.9846
Epoch [ 13/ 100] | d_loss: 1.2522 | g_loss: 0.9096
Epoch [ 13/ 100] | d_loss: 1.3038 | g_loss: 0.9086
Epoch [ 13/ 100] | d_loss: 1.3410 | g_loss: 0.9007
Epoch [ 14/ 100] | d_loss: 1.3972 | g_loss: 0.9263

```

Epoch [14/	100]		d_loss:	1.1457		g_loss:	0.9268
Epoch [14/	100]		d_loss:	1.4067		g_loss:	1.0591
Epoch [15/	100]		d_loss:	1.3855		g_loss:	1.0274
Epoch [15/	100]		d_loss:	1.2194		g_loss:	1.1702
Epoch [15/	100]		d_loss:	1.3796		g_loss:	0.8596
Epoch [16/	100]		d_loss:	1.2803		g_loss:	1.0645
Epoch [16/	100]		d_loss:	1.3096		g_loss:	0.9516
Epoch [16/	100]		d_loss:	1.3721		g_loss:	0.9343
Epoch [17/	100]		d_loss:	1.2190		g_loss:	1.2036
Epoch [17/	100]		d_loss:	1.2734		g_loss:	0.9071
Epoch [17/	100]		d_loss:	1.4173		g_loss:	0.9956
Epoch [18/	100]		d_loss:	1.2578		g_loss:	1.1108
Epoch [18/	100]		d_loss:	1.3411		g_loss:	0.9701
Epoch [18/	100]		d_loss:	1.4021		g_loss:	0.9006
Epoch [19/	100]		d_loss:	1.3386		g_loss:	0.9160
Epoch [19/	100]		d_loss:	1.2963		g_loss:	0.9935
Epoch [19/	100]		d_loss:	1.4204		g_loss:	0.9611
Epoch [20/	100]		d_loss:	1.3744		g_loss:	0.9034
Epoch [20/	100]		d_loss:	1.2739		g_loss:	0.9600
Epoch [20/	100]		d_loss:	1.3173		g_loss:	0.9407
Epoch [21/	100]		d_loss:	1.3409		g_loss:	0.9330
Epoch [21/	100]		d_loss:	1.3033		g_loss:	0.8945
Epoch [21/	100]		d_loss:	1.3469		g_loss:	0.9986
Epoch [22/	100]		d_loss:	1.3109		g_loss:	1.1370
Epoch [22/	100]		d_loss:	1.3431		g_loss:	1.1827
Epoch [22/	100]		d_loss:	1.4242		g_loss:	0.8984
Epoch [23/	100]		d_loss:	1.3188		g_loss:	1.0079
Epoch [23/	100]		d_loss:	1.2875		g_loss:	0.9586
Epoch [23/	100]		d_loss:	1.2436		g_loss:	1.0532
Epoch [24/	100]		d_loss:	1.3047		g_loss:	0.8399
Epoch [24/	100]		d_loss:	1.2922		g_loss:	1.0547
Epoch [24/	100]		d_loss:	1.3415		g_loss:	0.9550
Epoch [25/	100]		d_loss:	1.3227		g_loss:	1.0760
Epoch [25/	100]		d_loss:	1.3114		g_loss:	1.0052
Epoch [25/	100]		d_loss:	1.3168		g_loss:	0.8248
Epoch [26/	100]		d_loss:	1.3288		g_loss:	1.0172
Epoch [26/	100]		d_loss:	1.2819		g_loss:	0.8772
Epoch [26/	100]		d_loss:	1.3899		g_loss:	0.9913
Epoch [27/	100]		d_loss:	1.3186		g_loss:	0.9003
Epoch [27/	100]		d_loss:	1.3118		g_loss:	0.9462
Epoch [27/	100]		d_loss:	1.3907		g_loss:	1.1392
Epoch [28/	100]		d_loss:	1.2426		g_loss:	0.9351
Epoch [28/	100]		d_loss:	1.3180		g_loss:	0.9030
Epoch [28/	100]		d_loss:	1.4644		g_loss:	1.1859
Epoch [29/	100]		d_loss:	1.3197		g_loss:	0.7891
Epoch [29/	100]		d_loss:	1.2137		g_loss:	1.0703
Epoch [29/	100]		d_loss:	1.2770		g_loss:	0.9848
Epoch [30/	100]		d_loss:	1.2925		g_loss:	0.8826

Epoch [30/	100]		d_loss:	1.2893		g_loss:	0.8774
Epoch [30/	100]		d_loss:	1.2956		g_loss:	1.0330
Epoch [31/	100]		d_loss:	1.2196		g_loss:	0.8091
Epoch [31/	100]		d_loss:	1.3566		g_loss:	0.9080
Epoch [31/	100]		d_loss:	1.2753		g_loss:	1.1016
Epoch [32/	100]		d_loss:	1.2560		g_loss:	0.9505
Epoch [32/	100]		d_loss:	1.2989		g_loss:	1.0042
Epoch [32/	100]		d_loss:	1.3643		g_loss:	0.8927
Epoch [33/	100]		d_loss:	1.2769		g_loss:	0.9737
Epoch [33/	100]		d_loss:	1.3769		g_loss:	0.8789
Epoch [33/	100]		d_loss:	1.3109		g_loss:	1.0953
Epoch [34/	100]		d_loss:	1.3719		g_loss:	0.9287
Epoch [34/	100]		d_loss:	1.3272		g_loss:	0.9782
Epoch [34/	100]		d_loss:	1.3579		g_loss:	0.9943
Epoch [35/	100]		d_loss:	1.2647		g_loss:	0.9844
Epoch [35/	100]		d_loss:	1.3131		g_loss:	0.9080
Epoch [35/	100]		d_loss:	1.3689		g_loss:	0.9350
Epoch [36/	100]		d_loss:	1.2335		g_loss:	1.1166
Epoch [36/	100]		d_loss:	1.1683		g_loss:	1.1547
Epoch [36/	100]		d_loss:	1.3894		g_loss:	0.9600
Epoch [37/	100]		d_loss:	1.4048		g_loss:	1.0444
Epoch [37/	100]		d_loss:	1.2630		g_loss:	0.8688
Epoch [37/	100]		d_loss:	1.3267		g_loss:	0.8951
Epoch [38/	100]		d_loss:	1.2910		g_loss:	0.8282
Epoch [38/	100]		d_loss:	1.2974		g_loss:	1.0162
Epoch [38/	100]		d_loss:	1.3941		g_loss:	0.9535
Epoch [39/	100]		d_loss:	1.1903		g_loss:	1.2965
Epoch [39/	100]		d_loss:	1.2944		g_loss:	0.8642
Epoch [39/	100]		d_loss:	1.3383		g_loss:	0.8559
Epoch [40/	100]		d_loss:	1.1910		g_loss:	1.4142
Epoch [40/	100]		d_loss:	1.2928		g_loss:	1.0042
Epoch [40/	100]		d_loss:	1.4476		g_loss:	0.8785
Epoch [41/	100]		d_loss:	1.2930		g_loss:	1.1515
Epoch [41/	100]		d_loss:	1.2203		g_loss:	1.0401
Epoch [41/	100]		d_loss:	1.3833		g_loss:	0.9253
Epoch [42/	100]		d_loss:	1.2915		g_loss:	1.1200
Epoch [42/	100]		d_loss:	1.2022		g_loss:	1.2603
Epoch [42/	100]		d_loss:	1.3558		g_loss:	1.0367
Epoch [43/	100]		d_loss:	1.3766		g_loss:	1.1104
Epoch [43/	100]		d_loss:	1.2289		g_loss:	0.9819
Epoch [43/	100]		d_loss:	1.3489		g_loss:	0.8969
Epoch [44/	100]		d_loss:	1.3483		g_loss:	0.8587
Epoch [44/	100]		d_loss:	1.2589		g_loss:	0.8566
Epoch [44/	100]		d_loss:	1.4368		g_loss:	0.9920
Epoch [45/	100]		d_loss:	1.3017		g_loss:	1.4002
Epoch [45/	100]		d_loss:	1.2612		g_loss:	0.8207
Epoch [45/	100]		d_loss:	1.2975		g_loss:	0.9893
Epoch [46/	100]		d_loss:	1.3232		g_loss:	1.0155

Epoch [46/	100]	d_loss: 1.2396	g_loss: 1.3825
Epoch [46/	100]	d_loss: 1.3655	g_loss: 0.9999
Epoch [47/	100]	d_loss: 1.2516	g_loss: 0.9896
Epoch [47/	100]	d_loss: 1.2728	g_loss: 0.9268
Epoch [47/	100]	d_loss: 1.4520	g_loss: 1.0270
Epoch [48/	100]	d_loss: 1.3270	g_loss: 1.0549
Epoch [48/	100]	d_loss: 1.1733	g_loss: 0.9841
Epoch [48/	100]	d_loss: 1.3413	g_loss: 0.8965
Epoch [49/	100]	d_loss: 1.2776	g_loss: 0.8877
Epoch [49/	100]	d_loss: 1.2019	g_loss: 1.1288
Epoch [49/	100]	d_loss: 1.3470	g_loss: 1.0216
Epoch [50/	100]	d_loss: 1.2505	g_loss: 0.9932
Epoch [50/	100]	d_loss: 1.3520	g_loss: 0.8499
Epoch [50/	100]	d_loss: 1.2858	g_loss: 0.9268
Epoch [51/	100]	d_loss: 1.3898	g_loss: 1.0027
Epoch [51/	100]	d_loss: 1.2875	g_loss: 0.9086
Epoch [51/	100]	d_loss: 1.2878	g_loss: 1.0496
Epoch [52/	100]	d_loss: 1.2401	g_loss: 0.8165
Epoch [52/	100]	d_loss: 1.2490	g_loss: 0.7978
Epoch [52/	100]	d_loss: 1.2082	g_loss: 0.9915
Epoch [53/	100]	d_loss: 1.2312	g_loss: 1.0420
Epoch [53/	100]	d_loss: 1.3087	g_loss: 1.2269
Epoch [53/	100]	d_loss: 1.2136	g_loss: 1.1641
Epoch [54/	100]	d_loss: 1.2707	g_loss: 0.8842
Epoch [54/	100]	d_loss: 1.2299	g_loss: 0.8746
Epoch [54/	100]	d_loss: 1.3665	g_loss: 0.9369
Epoch [55/	100]	d_loss: 1.3157	g_loss: 0.8254
Epoch [55/	100]	d_loss: 1.1229	g_loss: 1.0649
Epoch [55/	100]	d_loss: 1.2802	g_loss: 0.9593
Epoch [56/	100]	d_loss: 1.5371	g_loss: 1.1198
Epoch [56/	100]	d_loss: 1.2728	g_loss: 0.9548
Epoch [56/	100]	d_loss: 1.3355	g_loss: 1.1381
Epoch [57/	100]	d_loss: 1.2697	g_loss: 0.9037
Epoch [57/	100]	d_loss: 1.2405	g_loss: 1.2255
Epoch [57/	100]	d_loss: 1.3482	g_loss: 1.1158
Epoch [58/	100]	d_loss: 1.3005	g_loss: 1.1478
Epoch [58/	100]	d_loss: 1.2590	g_loss: 1.0317
Epoch [58/	100]	d_loss: 1.2784	g_loss: 1.0695
Epoch [59/	100]	d_loss: 1.2271	g_loss: 0.9009
Epoch [59/	100]	d_loss: 1.2961	g_loss: 0.9430
Epoch [59/	100]	d_loss: 1.2620	g_loss: 0.9433
Epoch [60/	100]	d_loss: 1.2619	g_loss: 1.0115
Epoch [60/	100]	d_loss: 1.2794	g_loss: 1.0872
Epoch [60/	100]	d_loss: 1.4345	g_loss: 0.9186
Epoch [61/	100]	d_loss: 1.2988	g_loss: 0.9782
Epoch [61/	100]	d_loss: 1.1591	g_loss: 1.2982
Epoch [61/	100]	d_loss: 1.4653	g_loss: 0.9955
Epoch [62/	100]	d_loss: 1.3233	g_loss: 0.8551

Epoch [62/	100]		d_loss:	1.4645		g_loss:	0.9856
Epoch [62/	100]		d_loss:	1.3310		g_loss:	0.9223
Epoch [63/	100]		d_loss:	1.2351		g_loss:	1.6538
Epoch [63/	100]		d_loss:	1.2489		g_loss:	1.0497
Epoch [63/	100]		d_loss:	1.2768		g_loss:	1.0256
Epoch [64/	100]		d_loss:	1.3026		g_loss:	1.1597
Epoch [64/	100]		d_loss:	1.3820		g_loss:	1.1499
Epoch [64/	100]		d_loss:	1.3198		g_loss:	1.0013
Epoch [65/	100]		d_loss:	1.2573		g_loss:	1.2529
Epoch [65/	100]		d_loss:	1.2858		g_loss:	1.0339
Epoch [65/	100]		d_loss:	1.3242		g_loss:	0.9070
Epoch [66/	100]		d_loss:	1.3821		g_loss:	1.1427
Epoch [66/	100]		d_loss:	1.3039		g_loss:	0.9871
Epoch [66/	100]		d_loss:	1.3826		g_loss:	1.0136
Epoch [67/	100]		d_loss:	1.2751		g_loss:	0.7706
Epoch [67/	100]		d_loss:	1.1972		g_loss:	1.0919
Epoch [67/	100]		d_loss:	1.4427		g_loss:	0.9262
Epoch [68/	100]		d_loss:	1.2587		g_loss:	0.9697
Epoch [68/	100]		d_loss:	1.1228		g_loss:	1.2282
Epoch [68/	100]		d_loss:	1.4037		g_loss:	0.9404
Epoch [69/	100]		d_loss:	1.3538		g_loss:	0.8863
Epoch [69/	100]		d_loss:	1.2865		g_loss:	0.8829
Epoch [69/	100]		d_loss:	1.3817		g_loss:	1.0767
Epoch [70/	100]		d_loss:	1.3271		g_loss:	0.9144
Epoch [70/	100]		d_loss:	1.2499		g_loss:	1.0674
Epoch [70/	100]		d_loss:	1.2603		g_loss:	0.9379
Epoch [71/	100]		d_loss:	1.2161		g_loss:	1.0424
Epoch [71/	100]		d_loss:	1.3128		g_loss:	0.9293
Epoch [71/	100]		d_loss:	1.4150		g_loss:	1.0259
Epoch [72/	100]		d_loss:	1.3225		g_loss:	0.8122
Epoch [72/	100]		d_loss:	1.3034		g_loss:	1.2024
Epoch [72/	100]		d_loss:	1.4149		g_loss:	1.3195
Epoch [73/	100]		d_loss:	1.3639		g_loss:	0.9724
Epoch [73/	100]		d_loss:	1.1890		g_loss:	1.1638
Epoch [73/	100]		d_loss:	1.3886		g_loss:	0.9566
Epoch [74/	100]		d_loss:	1.3317		g_loss:	0.8894
Epoch [74/	100]		d_loss:	1.0961		g_loss:	1.2001
Epoch [74/	100]		d_loss:	1.3188		g_loss:	0.9089
Epoch [75/	100]		d_loss:	1.2584		g_loss:	1.0603
Epoch [75/	100]		d_loss:	1.2356		g_loss:	1.0055
Epoch [75/	100]		d_loss:	1.2381		g_loss:	0.9942
Epoch [76/	100]		d_loss:	1.3108		g_loss:	0.8810
Epoch [76/	100]		d_loss:	1.2471		g_loss:	0.9351
Epoch [76/	100]		d_loss:	1.3371		g_loss:	1.0227
Epoch [77/	100]		d_loss:	1.3044		g_loss:	1.0448
Epoch [77/	100]		d_loss:	1.3171		g_loss:	1.0235
Epoch [77/	100]		d_loss:	1.4401		g_loss:	0.9617
Epoch [78/	100]		d_loss:	1.2551		g_loss:	1.0200

Epoch [78/	100]	d_loss: 1.2172	g_loss: 0.9833
Epoch [78/	100]	d_loss: 1.4484	g_loss: 1.0597
Epoch [79/	100]	d_loss: 1.2945	g_loss: 1.0672
Epoch [79/	100]	d_loss: 1.2892	g_loss: 1.0825
Epoch [79/	100]	d_loss: 1.3873	g_loss: 1.1598
Epoch [80/	100]	d_loss: 1.3170	g_loss: 1.0550
Epoch [80/	100]	d_loss: 1.2631	g_loss: 0.9427
Epoch [80/	100]	d_loss: 1.3912	g_loss: 1.0001
Epoch [81/	100]	d_loss: 1.2396	g_loss: 1.2220
Epoch [81/	100]	d_loss: 1.2506	g_loss: 0.9819
Epoch [81/	100]	d_loss: 1.3788	g_loss: 0.9197
Epoch [82/	100]	d_loss: 1.3445	g_loss: 0.8363
Epoch [82/	100]	d_loss: 1.2379	g_loss: 1.1038
Epoch [82/	100]	d_loss: 1.3953	g_loss: 1.3281
Epoch [83/	100]	d_loss: 1.2572	g_loss: 1.1708
Epoch [83/	100]	d_loss: 1.2816	g_loss: 0.9025
Epoch [83/	100]	d_loss: 1.3750	g_loss: 0.8586
Epoch [84/	100]	d_loss: 1.3142	g_loss: 0.8887
Epoch [84/	100]	d_loss: 1.1706	g_loss: 1.0228
Epoch [84/	100]	d_loss: 1.2826	g_loss: 1.4092
Epoch [85/	100]	d_loss: 1.2179	g_loss: 1.1822
Epoch [85/	100]	d_loss: 1.2414	g_loss: 1.0184
Epoch [85/	100]	d_loss: 1.2499	g_loss: 0.9643
Epoch [86/	100]	d_loss: 1.4081	g_loss: 1.0215
Epoch [86/	100]	d_loss: 1.0672	g_loss: 1.2303
Epoch [86/	100]	d_loss: 1.3982	g_loss: 0.9861
Epoch [87/	100]	d_loss: 1.2966	g_loss: 0.9585
Epoch [87/	100]	d_loss: 1.2863	g_loss: 0.9498
Epoch [87/	100]	d_loss: 1.3952	g_loss: 0.9766
Epoch [88/	100]	d_loss: 1.3173	g_loss: 1.2159
Epoch [88/	100]	d_loss: 1.2444	g_loss: 1.1231
Epoch [88/	100]	d_loss: 1.3617	g_loss: 0.9537
Epoch [89/	100]	d_loss: 1.2070	g_loss: 1.0308
Epoch [89/	100]	d_loss: 1.3292	g_loss: 1.0234
Epoch [89/	100]	d_loss: 1.2550	g_loss: 1.0707
Epoch [90/	100]	d_loss: 1.2781	g_loss: 0.9172
Epoch [90/	100]	d_loss: 1.2674	g_loss: 1.0166
Epoch [90/	100]	d_loss: 1.3347	g_loss: 1.1140
Epoch [91/	100]	d_loss: 1.3387	g_loss: 0.9324
Epoch [91/	100]	d_loss: 1.2619	g_loss: 0.9601
Epoch [91/	100]	d_loss: 1.4142	g_loss: 1.0002
Epoch [92/	100]	d_loss: 1.3123	g_loss: 0.8796
Epoch [92/	100]	d_loss: 1.2449	g_loss: 0.9165
Epoch [92/	100]	d_loss: 1.2785	g_loss: 0.8928
Epoch [93/	100]	d_loss: 1.2613	g_loss: 1.4571
Epoch [93/	100]	d_loss: 1.2500	g_loss: 0.9886
Epoch [93/	100]	d_loss: 1.3941	g_loss: 0.9039
Epoch [94/	100]	d_loss: 1.2658	g_loss: 0.9556

```

Epoch [ 94/ 100] | d_loss: 1.2096 | g_loss: 1.0293
Epoch [ 94/ 100] | d_loss: 1.2934 | g_loss: 0.9862
Epoch [ 95/ 100] | d_loss: 1.3609 | g_loss: 1.4349
Epoch [ 95/ 100] | d_loss: 1.2245 | g_loss: 1.1067
Epoch [ 95/ 100] | d_loss: 1.3066 | g_loss: 0.9759
Epoch [ 96/ 100] | d_loss: 1.2949 | g_loss: 0.8319
Epoch [ 96/ 100] | d_loss: 1.3333 | g_loss: 0.9025
Epoch [ 96/ 100] | d_loss: 1.2014 | g_loss: 1.1053
Epoch [ 97/ 100] | d_loss: 1.2985 | g_loss: 1.0141
Epoch [ 97/ 100] | d_loss: 1.2627 | g_loss: 1.0140
Epoch [ 97/ 100] | d_loss: 1.3135 | g_loss: 0.9844
Epoch [ 98/ 100] | d_loss: 1.3431 | g_loss: 0.9857
Epoch [ 98/ 100] | d_loss: 1.1359 | g_loss: 0.8790
Epoch [ 98/ 100] | d_loss: 1.2968 | g_loss: 1.0595
Epoch [ 99/ 100] | d_loss: 1.2918 | g_loss: 1.0832
Epoch [ 99/ 100] | d_loss: 1.2491 | g_loss: 1.1004
Epoch [ 99/ 100] | d_loss: 1.3484 | g_loss: 0.9297
Epoch [ 100/ 100] | d_loss: 1.2533 | g_loss: 0.9658
Epoch [ 100/ 100] | d_loss: 1.2838 | g_loss: 1.0807
Epoch [ 100/ 100] | d_loss: 1.3661 | g_loss: 1.1609

```

2.8 Training loss

Here we'll plot the training losses for the generator and discriminator, recorded after each epoch.

```

In [44]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator')
         plt.plot(losses.T[1], label='Generator')
         plt.title("Training Losses")
         plt.legend()

```

```

Out[44]: <matplotlib.legend.Legend at 0x7fa4c8702160>

```




2.9 Generator samples from training


Here we can view samples of images from the generator. First we'll look at the images we saved during training.

```
In [45]: # helper function for viewing a list of passed in sample images
def view_samples(epoch, samples):
    fig, axes = plt.subplots(figsize=(7,7), nrows=4, ncols=4, sharey=True, sharex=True)
    for ax, img in zip(axes.flatten(), samples[epoch]):
        img = img.detach()
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        im = ax.imshow(img.reshape((28,28)), cmap='Greys_r')

In [46]: # Load samples from generator, taken while training
with open('train_samples.pkl', 'rb') as f:
    samples = pickle.load(f)
```

These are samples from the final training epoch. You can see the generator is able to reproduce numbers like 1, 7, 3, 2. Since this is just a sample, it isn't representative of the full range of images this generator can make.


```
In [47]: # -1 indicates final epoch's samples (the last in the list)
view_samples(-1, samples)
```



Below I'm showing the generated images as the network was training, every 10 epochs.

```
In [48]: rows = 10 # split epochs into 10, so 100/10 = every 10 epochs
        cols = 6
        fig, axes = plt.subplots(figsize=(7,12), nrows=rows, ncols=cols, sharex=True, sharey=True)

        for sample, ax_row in zip(samples[:, :int(len(samples)/rows)], axes):
            for img, ax in zip(sample[:, :int(len(sample)/cols)], ax_row):
                img = img.detach()
                ax.imshow(img.reshape((28,28)), cmap='Greys_r')
                ax.xaxis.set_visible(False)
                ax.yaxis.set_visible(False)
```



It starts out as all noise. Then it learns to make only the center white and the rest black. You can start to see some number like structures appear out of the noise like 1s and 9s.


2.10 Sampling from the generator

We can also get completely new images from the generator by using the checkpoint we saved after training. **We just need to pass in a new latent vector z and we'll get new samples!**

```
In [49]: # randomly generated, new latent vectors
sample_size=16
rand_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
rand_z = torch.from_numpy(rand_z).float()

G.eval() # eval mode
# generated samples
rand_images = G(rand_z)

# 0 indicates the first set of samples in the passed in list
# and we only have one batch of samples, here
view_samples(0, [rand_images])
```



In []:

In []: